Evaluation report

Human API · 26 October 2026 · synthetic concept analysis

1 · Scope and method

This report records checks on the local demo, using fictional people and files. No real employee data or Microsoft tenant was used. The demo uses word matching and fixed answer templates, not a language model, so these results say nothing about the quality or cost of a language model.

The backend tests check the question workflow and access rules. A separate exercise uses 36 written example questions to check whether the search finds the expected accessible file or returns no file when evidence is absent. It is a small word-matching check—not a test of whether an answer is truly correct, whether search understands paraphrases, or whether a production system is secure.

Run the checks with `make test` and `make eval`. Latest local run: 17 backend tests passed; all 36 search examples passed. The web interface also passed its type check and production build.

Example: the question about the 120-second timeout should find ADR-024. The word-matching check confirms that the file was found; it does not prove that the source fully answers whether 120 seconds is still needed.

2 · Observed results

Check	Result	What it means
Backend workflow tests	17 passed	The tested local rules behaved as expected
Written search examples	36/36 passed	32 expected files found; 4 no-evidence questions returned none
File-access violations	0 in these tests	Does not prove a real system is secure
Response time, cost, accepted answers	Not measured	No live service or user pilot

Behavior covered by tests

- Supported partial answer includes ADR-024 citation and limitation; restricted budget request denies without title/value.
- Low evidence escalates to backup; conflicting timeout documents require clarification.
- Human answer before deadline cancels pending job; retry limits stop repeated failure.
- Correction remains pending until owner approves; revocation suppresses previously returned answer and citations.
- Prompt injection fixture is treated as untrusted text; Restricted labels require explicit ACL.

3 · Interpretation and next evaluation

The 100% authored retrieval result is expected for a small synthetic vocabulary and should not be read as general accuracy. Lexical overlap misses paraphrases and can retrieve irrelevant material. Four denial cases do not establish robust resistance to inference or side channels. Prompt-injection sanitation uses a short marker list and is not a complete defense.

Before a pilot, measure

- Expert-labeled evidence recall and citation precision on representative, non-sensitive questions.
- Claim entailment, unsupported-answer rate, abstention precision and contradiction handling with blinded reviewers.
- Access regression across users, groups, sensitivity labels, revocation, deletion, tenant boundaries and cached responses.
- Concurrency, restart/recovery, notification idempotency, Graph throttling and unavailable service behavior.
- Active work time before/after, accepted resolution time, interruption reduction, escalation/correction rates and cost per accepted result.

Pilot safety gate: zero unauthorized disclosure. Stop and remediate on any access violation; do not trade that gate against answer coverage.